

RAGHAV KACHROO

(858)-241-1760 | rkachroo@ucsd.edu | [linkedin.com/in/raghavkachroo](https://www.linkedin.com/in/raghavkachroo) | github.com/mister-raggs

Software engineer drawn to friction points where systems are slow, manual, or breaking under load - building high-throughput data systems and ML inference infrastructure that fixes them at the root.

EDUCATION

University of California, San Diego Master of Science in Data Science (Artificial Intelligence & Machine Learning)	Sep 2024 – Mar 2026
Indraprastha Institute of Information Technology, Delhi Post Graduate Diploma in Data Science & Artificial Intelligence	Sep 2022 – Sep 2023
University of Delhi Bachelor of Management Studies	Jul 2018 – Jun 2021

EXPERIENCE

Hao AI Lab, UC San Diego Student Researcher	Jan 2026 – Present San Diego, CA
- Built a model-agnostic adaptive caching module for video-diffusion inference (residual-skip heuristic, no per-model adapters) delivering 22-47% latency reduction on Wan2.1 across an SSIM 0.875-0.946 quality frontier. - Surfaced a 23.7% <code>torch.compile</code> speedup on Wan2.1-1.3B via Inductor fusion (765 → 18 LayerNorm kernels); shipped FlashAttention as a <code>torch.library</code> custom op to extend the same graph-capture path to attention. - Ported FastVideo to NVIDIA's DGX Spark (GB10/Blackwell): enabled 4-bit (FP4) inference (-24% LTX2 denoise at 1080p) and shipped a sampling-preset fix (sharpness 92→431) upstream. - Cut Cosmos 2.5 inference latency 15.6% on A100 by quantizing decoded frames on-GPU before device transfer.	
Amazon Software Development Engineer Intern	Jun 2025 – Sep 2025 Bellevue, WA
- Built a distributed log indexing and query service over 42M+ log entries/hour using parallelized binary search, reducing incident triage latency from 15+ minutes to under 45 seconds. - Automated on-call SOPs using AWS Step Functions and Lambda, saving 12+ engineer-hours per week. - Exposed the log query service via MCP with caching and query batching, holding sub-2s response under load.	
Aark Global Software Developer, AI/ML	Apr 2023 – Sep 2024 Delhi, India
- Built an async document ingestion pipeline processing 18,000+ pages/day via Azure Queue Storage worker pool, maintaining throughput under variable load. - Implemented read/write routing during datastore migration - Cosmos DB reads, MongoDB writes - maintaining sub-100ms P95 latency throughout cutover. - Built a full-text search pipeline ingesting scanned PDFs through OCR into indexed Elasticsearch documents, enabling sub-180ms query latency over previously unsearchable content.	

PROJECTS & RESEARCH

Flare: ML-Powered Anomaly Detection Pipeline GitHub	Jan 2026 – Present
- Built an end-to-end ML anomaly detection pipeline applying Drain3 template mining, Isolation Forest scoring, and DBSCAN clustering to 11.2M log lines across 575K blocks, achieving 0.642 F1 at ~40ms CPU-only latency. - Built an LLM-as-judge eval harness grading incident summaries on relevance, specificity, and actionability, scoring 4.67/5 mean quality.	
log-triage-v2: High-Throughput Log Indexing Engine GitHub	Jan 2026 – Apr 2026
Go log indexing engine with time-sorted in-memory index, concurrent write ingestion, and binary search reads - sustains 42M rows/hour with sub-millisecond nearest-timestamp lookup.	
OneBusAway / maglev: OSS Contributions GitHub	Jul 2026 – Present
- Filed upstream REST API fix, removed off-spec pagination and corrected <code>includeReferences</code> handling. - Merged a <code>StaleDetector</code> fix correcting missed staleness detection in GTFS-realtime feed monitoring.	
TRAI: AI Mobile App to Reduce the Gap Between Triage and Care IEEE Publication	Jul 2025

TECHNICAL SKILLS

Languages: Python, Go, SQL, Java
Backend & Infra: FastAPI, REST APIs, Docker, Kubernetes, CI/CD, GitHub Actions, Linux, MCP
Distributed & Data Systems: Kafka, Spark, Airflow, Elasticsearch, MongoDB, Cosmos DB, Azure Queue Storage, Redis
Observability & Cloud: Prometheus, OpenTelemetry, AWS (Lambda, Step Functions, S3), Azure
AI & ML: PyTorch, Scikit-Learn, Hugging Face Transformers, LangChain, MLflow, RAG Pipelines, Diffusers, LoRA
Concepts: Distributed Systems, Concurrency, MLOps, Inference Optimization, System Design, Model Monitoring